



Using the IARC nextflow bioinformatics pipelines Day 2

Matthieu Foll
IARC course: May 23-24 2018

International Agency for Research on Cancer
Lyon, France

Agenda

Wednesday 23 May

09:00-10:00	Introduction to bioinformatics pipelines, nextflow, docker, Github and the IARC organization
10:00-10:30	Practical application: running your first pipeline
10:30-11:00	Break
11:00-11:30	The hidden structure of nextflow: work folder and configuration
11:30-12:30	Practical application: configuring, crashing, resuming and debugging pipelines

Thursday 24 May

09:00-09:30	Introduction to HPC clusters and running pipelines on a cluster.
09:30-10:30	Practical application: trace and visualise pipeline execution with log files.
10:30-11:00	Break
11:00-11:30	Introduction to the nextflow language: understanding what the pipelines are doing
11:30-12:30	Practical application: advanced usages toward reproducibility (choosing a container, Github releases and branches, modifying a pipeline etc.)



IARC bioinformatics platform

International Agency for Research on Cancer bioinformatics platform

Lyon, France <http://www.iarc.fr> folim@iarc.fr

Repositories 42 **People** 11 **Teams** 2 **Projects** 0 **Settings**

Pinned repositories

[Customize pinned repositories](#)

IARC-nf

List of IARC bioinformatics nextflow pipelines

★ 4 🍴 2

needlestack

Multi-sample somatic variant caller

● R ★ 16 🍴 4

mutspec

Mutation Spectra Analysis

● Perl ★ 5

alignment-nf

Whole Exome/Whole Genome Sequencing alignment pipeline

● Groovy ★ 4

RNAseq-nf

RNAseq analysis pipeline

● HTML ★ 2 🍴 1

mutect-nf

Mutect pipeline with Nextflow

● HTML ★ 3 🍴 1

Search repositories...

Type: All ▾

Language: All ▾

New

needlestack

Multi-sample somatic variant caller

[docker](#) [bam-files](#) [pipeline](#) [nextflow](#) [ngs](#) [regression](#)

● R ★ 16 🍴 4 Updated 5 days ago



RNAseq_analysis_scripts

Scripts for RNA seq analysis

● R ★ 1 Updated 14 days ago



Top languages

● R ● HTML ● Perl ● Shell
● Groovy

Most used topics

[Manage](#)

[nextflow](#) [ngs](#) [pipeline](#)
[bam-files](#) [alignment](#)

Our philosophy

- “Do It Once, Do It Right, And Use It Everywhere”
- “Keep it simple, stupid” (KISS principle):
 - most systems work best if they are kept simple
 - simplicity should be a key goal in design
 - code easier to maintain and to understand



Our design

- Too much automation is not for us:
 - Hard to read, to maintain and to keep modular
 - eg: we prefer to have one alignment pipeline; one variant calling pipeline; one annotation pipeline; one QC pipeline.
- One pipeline = one GitHub repository
- Docker containers hosted on DockerHub, compatible with Singularity
- CircleCI for tests and deployment
- Standardised readme, params, help etc. (one shared template)
- Use GitHub issues and releases
- Master branch ← beta branch ← dev branch

In practice

- Entry point: GitHub group
 - <https://github.com/IARCBioinfo>
- One central repo references all nextflow pipelines:
 - <https://github.com/IARCBioinfo/IARC-nf>
 - List pipelines with a short description
 - One pipeline = one repo, ends with “-nf”
 - Common instructions to use the pipelines (install nextflow, configuration, basic usage, docker...)
- A “template-nf” nextflow “hello-world” repo

High Performance Computing (HPC) cluster

Users



Management node



Network switch




Storage nodes



Backup



Computing nodes



International Agency for Research on Cancer

Job scheduler

- Computer program for controlling unattended background program execution of jobs
- Choose host based on:
 - Compute resource availability
 - Execution time allocated to user
 - Number of simultaneous jobs allowed for a user
 - Job priority
 - ...
- SGE, SLURM, PBS, Torque, **LSF** etc.

LSF most useful commands

- **bsub** command to launch a job. Useful options:
 - **-I**: Interactive job
 - **-oo**: output file name (overwrite)
 - **-eo**: error file name (overwrite)
 - **-J**: job name
 - **-m**: choose the hosts (or list of hosts e.g. `bsub -m "cn08 cn09 cn10"`)
 - **-n**: number of CPUs
 - **-R "rusage[mem=XXX]"** or **-r XXX**: ask for XXX MB of RAM
- **bjobs** to monitor jobs. Useful options:
 - **-w** for the full name
 - **-l** option for more details
 - **-u username** for a given user (or **-u all** for all users)

Nextflow is doing this for you!!!

Example

`samtools view -H NA06984.bam`

with docker becomes:

```
docker run -it --rm -v $PWD:$PWD -w $PWD --entrypoint /bin/bash samtools_img -c  
"samtools view -H NA06984.bam"
```

with LSF becomes:

```
bsub -oo out.txt -eo err.txt -n 1 -R "usage[mem=1000]" -M 1000 -J samtools  
"samtools view -H NA06984.bam"
```

with docker and LSF becomes:

```
bsub -oo out.txt -eo err.txt -n 1 -R "usage[mem=1000]" -M 1000 -J samtools  
"docker run -it --rm -v $PWD:$PWD -w $PWD --entrypoint /bin/bash samtools_img -c  
\"samtools view -H NA06984.bam\""
```

Hum... but we use singularity on our cluster... The syntax is actually different!
(singularity exec ...)

Nextflow is doing this for you!!!

Running on a cluster

```
nf_coverage_demo
~/nf_coverage_demo -- -bash
[x140083:nf_coverage_demo follm$ ssh follm@jupiter.iarc.fr
[follm@jupiter.iarc.fr's password:
Last login: Fri Mar  3 11:25:06 2017 from 10.10.13.47
[[follm@hn ~]$ cd /data/follm/nf_coverage_demo/
[[follm@hn nf_coverage_demo]$ nextflow run plot_coverage.nf --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Launching `plot_coverage.nf` [clever_ramanujan] - revision: 66d7be595f
[warm up] executor > local
[1e/ce2c42] Submitted process > coverage (3)
[f3/a818f6] Submitted process > coverage (1)
[2a/31fbc9] Submitted process > coverage (2)
[dc/ff61ac] Submitted process > coverage (6)
```

Running on the head-node: don't do that!!!

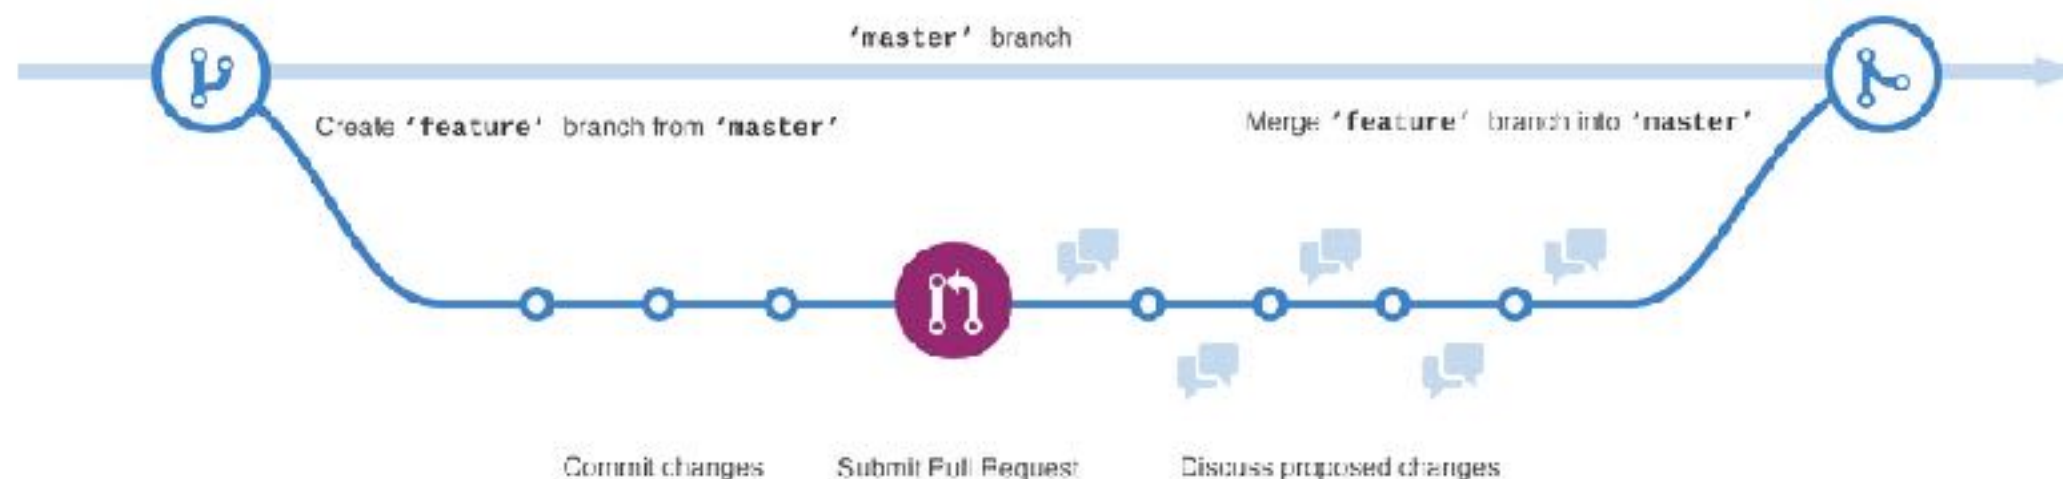
```
sftp://10.10.156.1//mnt/beegfs/follm/nf_coverage_demo/nextflow.config
1  process.executor='lsf'
2
```

```
[[follm@hn nf_coverage_demo]$ echo "process.executor='lsf'" > nextflow.config
[[follm@hn nf_coverage_demo]$ nextflow run plot_coverage.nf --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Launching `plot_coverage.nf` [gigantic_hawking] - revision: 66d7be595f
[warm up] executor > lsf
[37/af473d] Submitted process > coverage (4)
[ee/1ce1f9] Submitted process > coverage (1)
[fe/84c2d2] Submitted process > coverage (6)
```

Control queue size with `-qs` option

Speaking GitHub language

- **Branches**: master branch considered definitive, use other branches to experiment.
- Changes are called **commit**: one can **push** or **pull** commits to repository or a branch. Commits have unique IDs (hash).
- **Pull requests**: ask someone to **merge** the commits you did into their branch.
- **Releases** are user friendly commit IDs with **tag** names
- **Issues** are used to discuss bugs, features etc.



Nextflow and GitHub

- Any branch, tag or commit ID can be used to specify the revision you want to execute using the `-r` option:
`nextflow run nextflow-io/hello -r mybranch`
`nextflow run nextflow-io/hello -r v1.1`
- Be careful with docker version used:
 - Docker also has **tags** and hash IDs
 - If well done, a particular pipeline version will point to the right container automatically using the config file
 - If not you can manually specify the docker tag using:
`nextflow run -with-docker user/repo:tag`

Social coding

The histogram would be much nicer in pink #1

[Edit](#)[New issue](#)[Open](#)

mfoll opened this issue a minute ago · 1 comment



mfoll commented a minute ago

Member



No description provided.



mfoll added the **enhancement** label a minute ago



mfoll self-assigned this a minute ago



mfoll commented 33 seconds ago

Member



Thanks for your feedback, this will be implemented in v1.1

Assignees



mfoll

Labels



enhancement

Projects

None yet

Milestone



No milestone

International Agency for Research on Cancer



mfall added a commit that referenced this issue a minute ago

changed hist color to pink ... c75e7e0

mfall closed this just now

changed hist color to pink Browse files

Asked in [#1](#)

master

mfall committed on GitHub 13 seconds ago 1 parent [e89aba0](#) commit [c75e7e0a678905530fc71a9f02bdf4446da42d82](#)

Showing 1 changed file with 1 addition and 1 deletion. Unified Split

2 plot_coverage.nf View

@@ -57,7 +57,7 @@ process plot {

57	...	57	...
58	#!/usr/bin/env Rscript	58	#!/usr/bin/env Rscript
59	pdf("coverage.pdf")	59	pdf("coverage.pdf")
60	- hist(read.table("all_average.txt")[,1])	60	+ hist(read.table("all_average.txt")[,1],col="pink")
61	dev.off()	61	dev.off()
62	...	62	...
63	}	63	}

Latest release

v1.1

c75e7e0

v1.1



mfall released this 2 minutes ago

Changed histogram color to pink

Downloads



[Source code \(zip\)](#)



[Source code \(tar.gz\)](#)

v1.0

e89aba0

v1.0



mfall released this 7 minutes ago · 1 commit to master since this release

Update nextflow.config

Downloads



[Source code \(zip\)](#)



[Source code \(tar.gz\)](#)


International Agency for



Running specific versions

```
nf_coverage_demo
~/nf_coverage_demo — -bash
[x140083:nf_coverage_demo follm$ nextflow run iarcbioinfo/nf_coverage_demo -latest --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Pulling iarcbioinfo/nf_coverage_demo ...
Fast-forward
Launching `iarcbioinfo/nf_coverage_demo` [tiny_panini] - revision: c75e7e0a67 [master]
[warm up] executor > local
[88/be7ca4] Submitted process > coverage (2)
[68/25f37e] Submitted process > coverage (7)
```

Histogram of read.table("all_average.txt")[, 1]



```
nf_coverage_demo
~/nf_coverage_demo — -bash
[x140083:nf_coverage_demo follm$ nextflow run iarcbioinfo/nf_coverage_demo -r v1.0 --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Launching `iarcbioinfo/nf_coverage_demo` [amazing_almeida] - revision: e89aba0637 [v1.0]
[warm up] executor > local
[8a/1ffa6b] Submitted process > coverage (8)
[37/221eee] Submitted process > coverage (3)
[d8/7199ee] Submitted process > coverage (5)
```

Nextflow outputs

- stdout (what you see on the screen)
- work folder
- log file: `.nextflow.log`
- history: `.nextflow/history`
- Optional:
 - trace
 - timeline
 - report
 - email notification

History and log

~/nf_coverage_demo/nextflow/history

```
1 2017-03-03 10:39:57 2s backstabbing_shockley OK dds42716cd16e9073b6958b5d76f4d70 7ccf3889-b998-4c34-a2d4-ba5235c3ac83 nextflow run plot_coverage_1.nf --bam_folder BAM/ --bed TP53.bed
2 2017-03-03 10:40:43 2.7s compassionate_bohr OK dds42716cd16e9073b6958b5d76f4d70 c921a9c9-71b5-4279-831d-835da53f8f5a nextflow run plot_coverage_1.nf --bam_folder BAM/ --bed TP53.bed
3 2017-03-03 10:55:59 3.5s reverent_jones OK 66d7be595f210c509f1629889098c5ad 316692d7-bebb-42e9-911c-b13d80624d95 nextflow run plot_coverage_2.nf --bam_folder BAM/ --bed TP53.bed
4 2017-03-03 10:56:11 3.2s modest_shirley OK 6f32c834025fe85243ce5a62cd8b4780 9705a819-05f8-46d8-be5c-2e5414e1f9a9 nextflow run plot_coverage_2.nf --bam_folder BAM/ --bed TP53.bed
5 2017-03-03 11:07:01 3.7s nauseous_base OK 66d7be595f210c509f1629889098c5ad 10b19e2d-0cf2-481d-9111-70fa27b1198a nextflow run plot_coverage.nf --bam_folder BAM/ --bed TP53.bed
6 2017-03-03 11:13:01 2.9s high_blackwell ERR 66d7be595f210c509f1629889098c5ad 37d8590c-f6b1-4743-ac52-7d183d6df602 nextflow run plot_coverage.nf --bam_folder BAM/ --bed TP53.bed
```

~/nf_coverage_demo/nextflow.log

```
1 Mar-03 11:16:51.945 [main] DEBUG nextflow.cli.Launcher - S: /usr/local/bin/nextflow run plot_coverage.nf --with-trace --with-timeline --bam_folder BAM/ --bed TP53.bed
2 Mar-03 11:16:52.039 [main] INFO nextflow.cli.CmdRun - N E X T F L O W ~ version 0.23.4
3 Mar-03 11:16:52.048 [main] INFO nextflow.cli.CmdRun - Launching 'plot_coverage.nf' [cranky_leavitt] - revision: 66d7be595f
4 Mar-03 11:16:52.318 [main] DEBUG nextflow.Session - Session uuid: a88b516d-6f44-4041-b9f2-82e80c84b17e
5 Mar-03 11:16:52.318 [main] DEBUG nextflow.Session - Run name: cranky_leavitt
6 Mar-03 11:16:52.320 [main] DEBUG nextflow.Session - Executor pool size: 8
7 Mar-03 11:16:52.338 [main] DEBUG nextflow.cli.CmdRun -
8 Version: 0.23.4 build 4170
9 Modified: 24-02-2017 09:38 UTC (10:38 CEST)
10 System: Mac OS X 10.12.1
11 Runtime: Groovy 2.4.7 on Java HotSpot(TM) 64-Bit Server VM 1.8.0_51-b16
12 Encoding: UTF-8 (UTF-8)
13 Process: 30522@x1400083.local [10.10.13.47]
14 CPUs: 8 - Mem: 16 GB (47.1 MB) - Swap: 1 GB (938.2 MB)
15 Mar-03 11:16:52.345 [main] DEBUG nextflow.Session - Work-dir: /Users/follm/nf_coverage_demo/work [Mac OS X]
16 Mar-03 11:16:52.345 [main] DEBUG nextflow.Session - Script base path does not exist or is not a directory: /Users/follm/nf_coverage_demo/bin
17 Mar-03 11:16:52.428 [main] DEBUG nextflow.Session - Session start invoked
18 Mar-03 11:16:52.434 [main] DEBUG nextflow.processor.TaskDispatcher - Dispatcher > start
19 Mar-03 11:16:52.435 [main] DEBUG nextflow.trace.TraceFileObserver - Flow starting -- trace file: /Users/follm/nf_coverage_demo/trace.txt
20 Mar-03 11:16:52.438 [main] DEBUG nextflow.script.ScriptRunner - > Script parsing
21 Mar-03 11:16:52.536 [main] DEBUG nextflow.script.ScriptRunner - > Launching execution
22 Mar-03 11:16:52.555 [main] DEBUG nextflow.Channel - files for syntax: glob; folder: BAM/; pattern: *.bam; options: null
23 Mar-03 11:16:52.629 [main] DEBUG nextflow.processor.ProcessFactory - << taskConfig executor: null
24 Mar-03 11:16:52.629 [main] DEBUG nextflow.processor.ProcessFactory - >> processorType: 'local'
25 Mar-03 11:16:52.634 [main] DEBUG nextflow.executor.Executor - Initializing executor: local
26 Mar-03 11:16:52.636 [main] INFO nextflow.executor.Executor - [warn up] executor > local
27 Mar-03 11:16:52.641 [main] DEBUG n.processor.LocalPollingMonitor - Creating local task monitor for executor 'local' > cpus=8; memory=16 GB; capacity=8; pollInterval=100ms; dumpInterval=5m
28 Mar-03 11:16:52.644 [main] DEBUG nextflow.processor.TaskDispatcher - Starting monitor: LocalPollingMonitor
29 Mar-03 11:16:52.644 [main] DEBUG n.processor.TaskPollingMonitor - >>> barrier register (monitor: local)
30 Mar-03 11:16:52.646 [main] DEBUG nextflow.executor.Executor - Invoke register for executor: local
31 Mar-03 11:16:52.694 [main] DEBUG nextflow.Session - >>> barrier register (process: coverage)
32 Mar-03 11:16:52.703 [main] DEBUG nextflow.processor.TaskProcessor - Creating operator > coverage -- maxForks: 8
33 Mar-03 11:16:52.722 [main] DEBUG nextflow.processor.ProcessFactory - << taskConfig executor: null
34 Mar-03 11:16:52.723 [main] DEBUG nextflow.processor.ProcessFactory - >> processorType: 'local'
35 Mar-03 11:16:52.723 [main] DEBUG nextflow.executor.Executor - Initializing executor: local
36 Mar-03 11:16:52.723 [main] DEBUG nextflow.Session - >>> barrier register (process: mean)
37 Mar-03 11:16:52.729 [main] DEBUG nextflow.processor.TaskProcessor - Creating operator > mean -- maxForks: 8
38 Mar-03 11:16:52.850 [main] DEBUG nextflow.processor.ProcessFactory - << taskConfig executor: null
39 Mar-03 11:16:52.850 [main] DEBUG nextflow.processor.ProcessFactory - >> processorType: 'local'
40 Mar-03 11:16:52.850 [main] DEBUG nextflow.executor.Executor - Initializing executor: local
41 Mar-03 11:16:52.851 [main] DEBUG nextflow.Session - >>> barrier register (process: plot)
42 Mar-03 11:16:52.864 [main] DEBUG nextflow.processor.TaskProcessor - Creating operator > plot -- maxForks: 8
43 Mar-03 11:16:52.865 [main] DEBUG nextflow.script.ScriptRunner - > Await termination
44 Mar-03 11:16:52.865 [main] DEBUG nextflow.Session - Session await
45 Mar-03 11:16:52.990 [Actor Thread 4] INFO nextflow.Session - [cd/081600] Submitted process > coverage (2)
46 Mar-03 11:16:52.990 [Actor Thread 5] INFO nextflow.Session - [5e/f753d5] Submitted process > coverage (3)
47 Mar-03 11:16:52.990 [Actor Thread 9] INFO nextflow.Session - [9c/0cda8b] Submitted process > coverage (7)
48 Mar-03 11:16:52.990 [Actor Thread 8] INFO nextflow.Session - [4d/8bd5f0] Submitted process > coverage (6)
49 Mar-03 11:16:52.991 [Actor Thread 3] INFO nextflow.Session - [31/bf2185] Submitted process > coverage (1)
50 Mar-03 11:16:52.992 [Actor Thread 10] INFO nextflow.Session - [d1/8886c1] Submitted process > coverage (4)
```


Report (-with-report)

Nextflow workflow report

[ecstatic_booth]

Workflow execution completed successfully!

Run times

Tue May 22 18:31:22 CEST 2018 - Tue May 22 18:33:27 CEST 2018 (completed 3 hours ago, duration: 2m 6s)

187 succeeded

Nextflow command

```
nextflow run iarcbioinfo/nf_coverage_demo -latest -r v1.2 -with-docker --bam_folder data_test/BAM/BAM_multiple/ --bed data_test/BED/TP53_exon2_11.bed -with-report
```

CPU-Hours	0.1
Launch directory	/Users/foltn
Work directory	/Users/foltn/work
Project directory	/Users/foltn/.nextflow/assets/iarcbioinfo/nf_coverage_demo
Script name	plot_coverage.nf
Script ID	1bfa2de3de08d52bac30fdde640551c8
Workflow session	12396c92-954e-4f90-87e6-35547d3f77ec
Workflow repository	https://github.com/IARCbioinfo/nf_coverage_demo.git , revision v1.2 (commit hash 19c019913af89b4b8f6d93d115b136ef23cd0d74)
Workflow profile	standard
Workflow container	iarcbioinfo/nf_coverage_demo
Container engine	docker
Nextflow version	version 0.29.1, build 4804 (10-05-2018 07:47 UTC)

Trace (-with-trace)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N
1	task_id	hash	native_id	name	status	exit	submit	duration	realtime	%cpu	rss	vmem	rchar	wchar
2	2	cd/081600	30542	coverage (2)	COMPLETED	0	16:53.0	419ms	225ms	11.40%	7.3 MB	4.7 GB	0	0
3	4	9e/0cda8b	30543	coverage (7)	COMPLETED	0	16:53.0	446ms	397ms	-	-	-	-	-
4	3	5e/f753d5	30546	coverage (3)	COMPLETED	0	16:53.0	459ms	411ms	-	-	-	-	-
5	6	dc/087bc6	30541	coverage (5)	COMPLETED	0	16:53.0	467ms	417ms	-	-	-	-	-
6	5	c6/83374d	30539	coverage (4)	COMPLETED	0	16:53.0	550ms	283ms	2.30%	6.6 MB	4.7 GB	0	0
7	1	31/bf2185	30544	coverage (1)	COMPLETED	0	16:53.0	641ms	591ms	-	-	-	-	-
8	8	d7/18814e	30545	coverage (8)	COMPLETED	0	16:53.0	695ms	646ms	-	-	-	-	-
9	11	f0/67de05	30949	coverage (11)	COMPLETED	0	16:53.4	240ms	227ms	-	-	-	-	-
10	14	2e/dce17f	30928	coverage (14)	COMPLETED	0	16:53.4	352ms	80ms	0.70%	1 MB	2.3 GB	0	0
11	9	32/05246f	31055	coverage (9)	COMPLETED	0	16:53.5	342ms	84ms	0.70%	1 MB	2.3 GB	0	0
12	16	2c/01e5ea	31195	coverage (16)	COMPLETED	0	16:53.7	276ms	127ms	0.70%	8.6 MB	4.7 GB	0	0
13	17	a5/8e5a2e	31288	coverage (17)	COMPLETED	0	16:53.8	231ms	127ms	-	-	-	-	-
14	12	ee/c77b04	31191	coverage (12)	COMPLETED	0	16:53.7	340ms	95ms	0.70%	1 MB	2.3 GB	0	0
15	18	5d/0f162d	31449	mean (1)	COMPLETED	0	16:53.9	158ms	74ms	-	-	-	-	-
16	15	72/ddf049	31146	coverage (15)	COMPLETED	0	16:53.6	420ms	219ms	10.80%	9.6 MB	4.7 GB	0	0
17	10	95/26c5e0	30954	coverage (10)	COMPLETED	0	16:53.5	634ms	542ms	-	-	-	-	-
18	19	65/d41f91	31517	mean (2)	COMPLETED	0	16:54.0	176ms	142ms	-	-	-	-	-
19	21	3b/800cd4	31566	mean (3)	COMPLETED	0	16:54.0	181ms	161ms	-	-	-	-	-
20	23	0e/e9abc8	31599	mean (4)	COMPLETED	0	16:54.1	173ms	127ms	-	-	-	-	-
21	20	4d/b7bb4e	31534	coverage (18)	COMPLETED	0	16:54.0	269ms	237ms	-	-	-	-	-
22	22	36/e28f53	31583	coverage (19)	COMPLETED	0	16:54.0	231ms	217ms	-	-	-	-	-
23	7	4d/8bd5f0	30540	coverage (6)	COMPLETED	0	16:53.0	1.4s	358ms	2.50%	5.2 MB	4.7 GB	0	0
24	28	ac/c45043	31852	mean (6)	COMPLETED	0	16:54.3	169ms	157ms	-	-	-	-	-
25	25	de/2ed75c	31715	coverage (21)	COMPLETED	0	16:54.1	346ms	188ms	10.60%	10.5 MB	4.7 GB	0	0
26	24	e0/450f34	31649	coverage (20)	COMPLETED	0	16:54.1	438ms	386ms	-	-	-	-	-
27	27	ff/c08e08	31807	coverage (22)	COMPLETED	0	16:54.2	307ms	153ms	0.00%	10.7 MB	4.7 GB	0	0
28	31	a0/337b22	32049	mean (8)	COMPLETED	0	16:54.4	295ms	242ms	-	-	-	-	-
29	34	d5/04b9a6	32117	mean (9)	COMPLETED	0	16:54.5	204ms	104ms	-	-	-	-	-
30	33	a1/b9754f	32116	coverage (25)	COMPLETED	0	16:54.5	246ms	144ms	-	-	-	-	-
31	13	1a/2bbf4f	30935	coverage (13)	COMPLETED	0	16:53.4	1.4s	1.2s	100.70%	2.7 MB	4.7 GB	0	0

Practical

- Run pipelines on the cluster using Singularity `-with-singularity`
- Explore all nextflow outputs/logs, including in the `work` directory
- Try to run without docker/singularity
- Try to run v1.0, still with docker/singularity
- Try different combinations of pipeline version and container version
- Look at the help of the `nextflow` commands to find useful options (in particular `run -latest, -qs, -bg`)
- Look at the `clone, drop, list` and `pull` commands to manage pipelines
- Try to modify a pipeline (feel free to send me a PR!)
- Check if a new version of nextflow itself is available using `nextflow self-update`

Dataflow programming

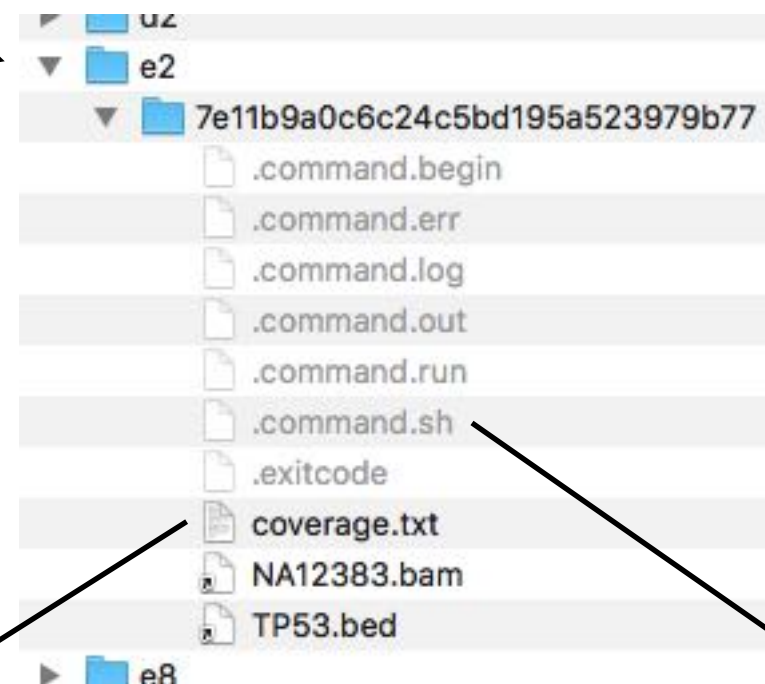
- Traditionally, a program is a series of operations happening in a specific order (“sequential programming”)
- Dataflow programming emphasizes the movement of data and models programs as a series of connections:
 - Explicitly defined inputs and outputs connect operations, which function like black boxes.
 - An operation runs as soon as all of its inputs become valid. Thus, dataflow languages are inherently parallel.

Our first process

```
1
2 // Defines pipeline parameters
3 params.bam_folder = null
4 params.bed = null
5
6 // The bed file
7 bed = file(params.bed)
8
9 // Creates the `bam` channel
10 bam = Channel.fromPath( params.bam_folder+'/*.bam' )
11
12 // Step 1. launch bedtools software to calculate coverage at each position of the bed
13 process coverage {
14
15     input:
16     file bam
17     file bed
18
19     output:
20     file 'coverage.txt' into coverage
21
22     shell:
23     '''
24     bedtools coverage -d -a !{bed} -b !{bam} > coverage.txt
25     '''
26 }
```


Running the pipeline

```
nf_coverage_demo — -bash — 129x33
[x140083:nf_coverage_demo follm$ nextflow run plot_coverage_1.nf --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Launching `plot_coverage_1.nf` [compassionate_bohr] - revision: dd542716cd
[warm up] executor > local
[e2/7e11b9] Submitted process > coverage (5)
[a4/2206c6] Submitted process > coverage (7)
[12/62bcf4] Submitted process > coverage (8)
[80/f6a03f] Submitted process > coverage (1)
[41/afb57a] Submitted process > coverage (3)
[95/425153] Submitted process > coverage (6)
[8f/44fcb4] Submitted process > coverage (2)
[e8/fe8a8f] Submitted process > coverage (4)
[2b/c8410d] Submitted process > coverage (14)
[30/9157d7] Submitted process > coverage (15)
[46/f2eab0] Submitted process > coverage (9)
[59/8e58d8] Submitted process > coverage (12)
[fd/c188c3] Submitted process > coverage (13)
[a9/444392] Submitted process > coverage (10)
[11/3b3c1f] Submitted process > coverage (11)
[ce/51d91c] Submitted process > coverage (16)
[4c/511af0] Submitted process > coverage (17)
```



```
~/nf_coverage_demo/work/e2/7e11b9a0c6c24c5bd195a523979b77/.command.sh
1  #!/bin/bash -ue
2  bedtools coverage -d -a TP53.bed -b NA12383.bam > coverage.txt
```

```
~/nf_coverage_demo/work/e2/7e11b9a0c6c24c5bd195a523979b77/coverage.txt
1  17  7579854  7579962  TP53_exon2  1  81
2  17  7579854  7579962  TP53_exon2  2  81
3  17  7579854  7579962  TP53_exon2  3  82
4  17  7579854  7579962  TP53_exon2  4  82
```

Adding a second process

```
19     output:
20     file 'coverage.txt' into coverage
21
22     shell:
23     '''
24     bedtools coverage -d -a !{bed} -b !{bam} > coverage.txt
25     '''
26 }
27
28 // Step 2. launch custom awk script to calculate the mean coverage
29 process mean {
30
31     input:
32     file coverage
33
34     output:
35     stdout average
36
37     shell:
38     '''
39     awk '{ sum += $6 } END { if (NR > 0) print sum / NR }' !{coverage}
40     '''
41 }
```

Running the pipeline

```
nf_coverage_demo -- -bash -- 129x33
[x140083:nf_coverage_demo follm$ nextflow run plot_coverage_2.nf --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Launching `plot_coverage_2.nf` [modest_shirley] - revision: 6f32c83402
[warm up] executor > local
[f6/cc4a86] Submitted process > coverage (7)
[e7/1c6c05] Submitted process > coverage (1)
[1f/08005d] Submitted process > coverage (4)
[37/0db07c] Submitted process > coverage (8)
[ba/7092ed] Submitted process > coverage (3)
[1b/e48edf] Submitted process > coverage (2)
[e3/c18377] Submitted process > coverage (12)
[89/fde6b3] Submitted process > coverage (5)
[d3/456b82] Submitted process > coverage (6)
[b6/59320b] Submitted process > coverage (11)
[c9/319f64] Submitted process > coverage (10)
[e5/be57ba] Submitted process > coverage (9)
[9a/ed4a42] Submitted process > coverage (13)
[28/243d0a] Submitted process > coverage (14)
[77/6d9dee] Submitted process > coverage (15)
[c4/2c85af] Submitted process > coverage (16)
[17/f2b669] Submitted process > coverage (17)
[a8/e2f0a9] Submitted process > coverage (18)
[87/7e3a9b] Submitted process > coverage (19)
[c5/7489dd] Submitted process > mean (1)
[25/8f2f62] Submitted process > mean (2)
[56/d40ff8] Submitted process > mean (3)
[be/02c2d1] Submitted process > mean (4)
[ce/0d34de] Submitted process > coverage (21)
```

```
~/nf_coverage_demo/work/c5/7489ddc158fe4aecc3eeeb4f6495e3/.command.sh
1  #!/bin/bash -ue
2  awk '{ sum += $6 } END { if (NR > 0) print sum / NR }' coverage.txt
3
```

▼ c5

▼ 7489ddc158fe4aecc3eeeb4f6495e3

- coverage.txt
- .exitcode
- .command.sh
- .command.run
- .command.out
- .command.log
- .command.err
- .command.begin

```
~/nf_coverage_demo/work/c5/7489ddc158fe4aecc3eeeb4f6495e3/.command.out
1  61.4442
2
```


The final process

```
34     output:
35     stdout average
36
37     shell:
38     '''
39     awk '{ sum += $6 } END { if (NR > 0) print sum / NR }' !{{coverage}}
40     '''
41 }
42 // collect output of all means to a single file
43 all_average = average.collectFile(name: 'all_average.txt')
44
45 // Step 3: plot histogram of mean coverage using a custom R script
46 process plot {
47
48     input:
49     file all_average
50
51     output:
52     file 'coverage.pdf'
53
54     publishDir '.', mode: 'move'
55
56     shell:
57     '''
58     #!/usr/bin/env Rscript
59     pdf("coverage.pdf")
60     hist(read.table("all_average.txt")[,1])
61     dev.off()
62     '''
63 }
```

This is our
final result

Running the pipeline

nf_coverage_demo — -bash — 129x33

```
[x140083:nf_coverage_demo follm$ nextflow run plot_coverage.nf --bam_folder BAM/ --bed TP53.bed
N E X T F L O W ~ version 0.23.4
Launching `plot_coverage.nf` [nauseous_bose] - revision: 66d7be595f
[warm up] executor > local
[33/2d903b] Submitted process > coverage (2)
[77/c7a53a] Submitted process > coverage (8)
[3e/576a19] Submitted process > coverage (7)
[d0/933d8d] Submitted process > coverage (6)
[28/5e52b6] Submitted process > coverage (5)
[7f/05adf5] Submitted process > coverage (1)
[ed/d4b830] Submitted process > coverage (4)
[42/109ee8] Submitted process > coverage (3)
[88/e83b10] Submitted process > coverage (11)
[0b/d20ce4] Submitted process > coverage (15)
[fe/6dfaf9] Submitted process > coverage (14)
[f3/fe6bd5] Submitted process > mean (3)

...

[7d/04a7b7] Submitted process > mean (37)
[6d/2b2c77] Submitted process > mean (38)
[8c/6eeada] Submitted process > mean (39)
[af/cc314a] Submitted process > mean (40)
[8e/e458cc] Submitted process > plot (1)
```

~/nf_coverage_demo/

1	81.2483
2	142.893
3	158.128
4	56.5989
5	61.7087
6	156.421
7	53.7202
8	55.1144
9	111.836
10	260.168
11	167.545
12	242.676
13	167.266
14	73.396
15	154.321
16	120.37

8e

e458cc80e51f0a8485907774b829e7


- all_average.txt
- .exitcode
- .command.sh
- .command.run
- .command.out
- .command.log
- .command.err
- .command.begin

~/nf_coverage_demo/work/8e/e458cc80e51f0a8485907774b829e7/.command.sh

```
1 #!/usr/bin/env Rscript
2 pdf("coverage.pdf")
3 hist(read.table("all_average.txt")[,1])
4 dev.off()
```

Result

Histogram of read.table("all_average.txt")[, 1]



```

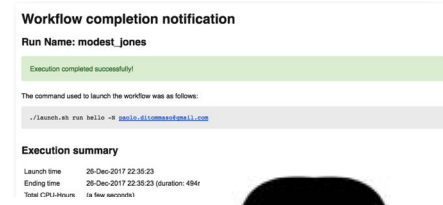
1  #!/usr/bin/env Rscript
2  // Defines pipeline parameters
3  params.bam_folder = null
4  params.bed = null
5
6  // The bed file
7  bed = file(params.bed)
8
9  // Creates the 'bam' channel
10 bam = Channel.fromPath( params.bam_folder+'/*.bam' )
11
12 // Step 1. launch bedtools software to calculate coverage at each position of the bed
13 process coverage {
14   inputs:
15     file bam
16     file bed
17
18   output:
19     file 'coverage.txt' into coverage
20
21   shell:
22     ...
23     bedtools coverage -d -a {bed} -b {bam} > coverage.txt
24     ...
25 }
26
27 // Step 2. launch custom awk script to calculate the mean coverage
28 process mean {
29   inputs:
30     file coverage
31
32   output:
33     stdout average
34
35   shell:
36     ...
37     awk '{ sum += $6 } END { if (NR > 0) print sum / NR }' {coverage}
38     ...
39 }
40
41 // collect output of all means to a single file
42 all_average = average.collectFile(name: 'all_average.txt')
43
44 // Step 3: plot histogram of mean coverage using a custom R script
45 process plot {
46   inputs:
47     file all_average
48
49   output:
50     file 'coverage.pdf'
51
52   publishDir '.', mode: 'move'
53
54   shell:
55     ...
56     #!/usr/bin/env Rscript
57     pdf("coverage.pdf")
58     hist(read.table("all_average.txt")[,1])
59     dev.off()
60     ...
61 }

```


Full life cycle



4 Nextflow runs each *process* of the pipeline inside the container on the available *executors* (workstation, HPC, cloud)



5 Nextflow sends an execution *report*



1 A **user** runs the pipeline using *nextflow*. If something goes wrong he opens an *issue*.

nextflow

3 Nextflow downloads the docker container with all the *software* from DockerHub



5 If circleci checks have passed, a new **docker** container is build and hosted on DockerHub


2 Nextflow downloads the pipeline from github



2 Github tracks the *changes* and creates a unique identifier for this *commit*

1 A **developer** modifies the pipeline and pushes the changes to github (*commit*).

This is the only manual action he has to perform as all the following steps are triggered automatically




3 Circleci starts running *tests* in the cloud to check if the changes haven't broken something and are still producing expected results on some test datasets

SUCCESS	master #44	7 months ago	02:59	6486dbc	1.0
FIXED	master #43	7 months ago	03:51	5cf614c	1.0
FAILED	master #42	7 months ago	02:42	f9d384b	1.0

4 Circleci sends feedback on the pipeline github webpage about the results of the test

circleci passing



Final word

Don't try to twist nextflow to fit your current habits,
give it a chance to actually change your habits